

E-jection fraction - Tracking how your website pumps out E-Mails

Ben Trovato^{*}
Institute for Clarity in
Documentation
1932 Wallamaloo Lane
Wallamaloo, New Zealand
trovato@corporation.com

G.K.M. Tobin[†]
Institute for Clarity in
Documentation
P.O. Box 1212
Dublin, Ohio 43017-6221
webmaster@marysville-
ohio.com

Lars Thørvæld[‡]
The Thørvæld Group
1 Thørvæld Circle
Hekla, Iceland
larst@affiliation.org

Lawrence P. Leipuner
Brookhaven Laboratories
Brookhaven National Lab
P.O. Box 5000
lleipuner@researchlabs.org

Sean Fogarty
NASA Ames Research Center
Moffett Field
California 94035
fogartys@amesres.org

Charles Palmer
Palmer Research Laboratories
8600 Datapoint Drive
San Antonio, Texas 78229
cpalmer@prl.com

ABSTRACT

E-Mail header injection vulnerability is a class of vulnerability that can occur in web applications that use user input to construct e-mail messages. E-Mail injection is possible when the mailing script fails to check for the presence of e-mail headers in user input (either form fields or URL parameters). The vulnerability exists in the reference implementation of the built-in “mail” functionality in popular languages like PHP, Java, Python, and Ruby. With the proper injection string, this vulnerability can be exploited to inject additional headers and/or modify existing headers in an E-Mail message, allowing an attacker to completely alter the content of the e-mail.

This thesis develops a scalable mechanism to automatically detect E-Mail Header Injection vulnerability and uses this mechanism to quantify the prevalence of EMail Header Injection vulnerabilities on the Internet. Using a black-box testing approach, the system crawled 21,675,680 URLs to find URLs which contained form fields. 6,794,917 such forms were found by the system, of which 1,132,157 forms contained e-mail fields. The system used this data feed to discern the forms that could be fuzzed with malicious payloads. Amongst the 934,016 forms tested, 52,724 forms were found to be injectable with more malicious payloads. The system tested 46,156 of these, and was able to find 496 vulnerable URLs across 222 domains, which proves that the threat is

widespread and deserves future research attention.

CCS Concepts

•Computer systems organization → Embedded systems; *Redundancy*; Robotics; •Networks → Network reliability;

Keywords

ACM proceedings; L^AT_EX; text tagging

1. INTRODUCTION

The World Wide Web has single-handedly brought about a change in the way we use computers. The ubiquitous nature of the Web has made it possible for the general public to access it anywhere and on multiple devices like phones, laptops, personal digital assistants, and even on TVs and cars. This has ushered in an era of responsive web applications which depend on user input. While this rapid pace of development has improved the speed of dissemination of information, it does come at a cost. Attackers have an added incentive to break into user’s e-mail accounts more than ever. E-Mail accounts are usually connected to almost all other online accounts of a user, and e-mails continue to serve as the principal mode of official communication on the web for most institutions. Thus, the impact an attacker can have by having control over the e-mail communication sent by websites to users is of an enormous magnitude.

Since attackers typically masquerade themselves as users of the system, if user input is to be trusted, then developers need to have proper sanitization routines in place. Many different injection attacks such as SQL injection or cross-site scripting (XSS) [13] are possible due to improper sanitization of user input.

Our research focuses on a lesser known injection attack known as E-Mail Header Injection. E-Mail Header Injection can be considered as the e-mail equivalent of HTTP Header Injection vulnerability [7]. The vulnerability exists in the reference implementation of the built-in “mail” functionality in popular languages like PHP, Java, Python, and Ruby.

^{*}Dr. Trovato insisted his name be first.

[†]The secretary disavows any knowledge of this author’s actions.

[‡]This author is the one who did all the really hard work.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

WOODSTOCK '97 El Paso, Texas USA

© 2016 ACM. ISBN 123-4567-24-567/08/06...\$15.00

DOI: 10.475/123_4

With the proper injection string, this vulnerability can be exploited to inject additional headers and/or modify existing headers in an e-mail message — with the potential to alter the contents of the e-mail message — while still appearing to be from a legitimate source.

E-Mail Header Injection attacks have the potential to allow an attacker to perform e-mail spoofing, resulting in phishing attacks that can lead to identity theft. The objective of our research is to study the prevalence of this vulnerability on the World Wide Web, and identify whether further research is required in this area.

We performed an expansive crawl of the web, extracting forms with e-mail fields, and injecting them with different payloads to infer the existence of E-Mail Header Injection vulnerability. We then audited received e-mails to see if any of the injected data was present. This allowed us to classify whether a particular URL was vulnerable to the attack. The entire system works in a black-box manner, without looking at the web application's source code, and only analyzes the e-mails we receive based on the injected payloads.

Structure of document

This thesis document is divided logically into the following sections:

- Chapter 2 discusses the background of E-Mail Header Injection, a brief history of the vulnerability, and enumerates the languages and platforms affected by this vulnerability.
- Chapter 3 discusses the System design, the architecture, and the components of the system.
- Chapter 4 describes the experimental setup and sheds light on how we overcame the issues and assumptions discussed in Chapter 3.
- Chapter 5 presents our findings and our analysis of the results.
- Chapter 6 continues the discussion of the results; the lessons learned over the course of the project, limitations, and a suitable mitigation strategy to overcome the vulnerability.
- Chapter 7 explores related work in the area.
- Chapter 8 concludes this thesis, with ideas to expand the research in this area.

We hope that our research sheds some light on this relatively less well-known vulnerability, and find out its prevalence on the World Wide Web. In summary, we make the following contributions:

- A black-box approach to detecting the presence of E-Mail Header Injection vulnerability in a web application.
- A detection and classification tool based on the above approach, which will automatically detect such E-Mail Header Injection vulnerabilities in a web application.
- A quantification of the presence of such vulnerabilities on the World Wide Web, based on a crawl of the Web, including 21,675,680 URLs and 6,794,917 forms.

2. E-MAIL HEADER INJECTION BACKGROUND

This chapter goes into the background of the problem at hand and gives a brief history of E-Mail Header Injection. It then describes the languages affected by this vulnerability and discusses the overall impact E-Mail Header Injection can have, and the attacks that can result from this vulnerability.

2.1 Problem Background

E-Mail Header Injection belongs to a broad class of vulnerabilities known simply as injection attacks. However, unlike its more popular siblings, SQL injection [2], [5], [17], Cross-Site Scripting (XSS) [8], [10] or even HTTP Header Injection [9], relatively little research is available on E-Mail Header Injection.

As with other vulnerabilities in this class, E-Mail Header Injection is caused due to improper sanitization (or lack thereof) of user input. If the script that constructs e-mails from user input fails to check for the presence of e-mail headers in the user input, a malicious user — using a well-crafted payload — can control the headers set for this particular e-mail. This can be leveraged to enable malicious attacks, including, but not limited to, spoofing, phishing, etc.

2.2 History of E-Mail Injection

E-Mail Header Injection seems to have been first documented over a decade ago, in a late 2004 article on phpsecure.info [20] accredited to user tobozo@phpsecure.info describing how this vulnerability existed in the reference implementation of the “mail” function in PHP, and how it can be exploited. More recently, a blog post by Damon Kohler [11] and an accompanying wiki article [4] describe the attack vector and outline a few defense measures for the same.

As this vulnerability was initially found in the *mail()* function of PHP, E-Mail Header Injection can be traced to as early as the beginning of the 2000's, present in the *mail()* implementation of PHP 4.0.

The vulnerability was also described briefly (less than a page) by Stuttard and Pinto in their widely acclaimed book, “*The Web Application Hacker's Handbook: Discovering and Exploiting Security Flaws*” [19]. A concise timeline of the vulnerability is presented in Table 1.

An example of the vulnerable code written in PHP is shown in Listing 1. This code takes in user input from the PHP superglobal “\$_REQUEST['email']”, and stores it in the variable “\$from”, which is later passed to the “mail” function to construct and send the e-mail.

```
1 $from = $_REQUEST['email'];
2 $subject = "Hello Sai Pc";
3 $message = "We need you to reset your
4   password";
5 $to = "schand31@asu.edu";
6 // example attack string to be injected as
7 // the value for
8 // $_REQUEST['email'] => 'sai@sai.com\nCC:
9   spc@spc.com'
10 $returnValue = mail($to, $subject, $message, "
11   From: $from");
12 // E-Mail gets sent to both schand31@asu.
13   edu AND spc@spc.com
```

Listing 1: PHP program with e-mail header injection vulnerability.

When this code is given the malicious input “sai@sai.com\nBCC:spc@spc.com” as the value of the “\$_REQUEST['email']”, it generates the SMTP Headers shown in Listing 2. It can be seen that the ‘CC’ (carbon copy) header that we injected appears as part of the resulting SMTP message. This will make the e-mail get sent to the e-mail address specified as part of the ‘CC’ as well.

```
1 Received: from mail.ourdomain.com
  ([62.121.130.29])
2   by sai.com (Postfix) with ESMTP id
   5A08E52C0154
3   for <sai@sai.com>; Sun, 20 Mar 2016
   13:56:58 -0700 (MST)
4 To: sai@sai.com
5 Subject: Hello Sai Pc
6 CC: spc@spc.com
7 Date: Sun, 20 Mar 2016 13:56:58 -0700 (MST)
8
9 We need you to reset your password
```

Listing 2: SMTP headers generated by a PHP mailing script.

Year	Notes
Early 2000's	PHP 4.0 is released, along with support for the tion against E-Mail Header Injection.
Jul 2004	Next Major version of PHP - Version 5.0 relea
Dec 2004	First known article about the vulnerability su
Oct 2007	The vulnerability makes its way into a text by
Dec 2008	Blog post and accompanying wiki about the examples.
Apr 2009	Bug filed about email.header package to fix th
Jan 2011	Bug fix for Python 3.1, Python 3.2, Python, 2 to older versions not available.
Sep 2011	The vulnerability is described with an exam Stuttard and Pinto.
Aug 2013	Acunetix adds E-Mail Header Injection to th part of their Enterprise Web Vulnerability Sca
May 2014	Security Advisory for JavaMail SMTP Heade written by Alexandre Herzog.
Dec 2015	PHP 7 releases, mail function still unpatched?

Table 1: A brief history of e-mail header injection.

2.3 Languages Affected

This section describes the popular languages which exhibit this type of vulnerability. This section is not intended as a complete reference of vulnerable functions and methods, but rather as a guide that specifies which parts of the language are known to have the vulnerability.

2.3.1 PHP

PHP was one of the first languages found to have this vulnerability in its implementation of the *mail()* function. The early finding of this vulnerability can be attributed in part to the success and popularity of the language for creating web pages. According to w3techs [22], PHP is used by 81.9% of all the websites in existence, thereby creating the possibility of this vulnerability to be widespread.

PHP's low barrier to entry and lack of developer education about the existence of this vulnerability have contributed to the vulnerability continuing to exist in the language. After

13 further iterations of the language since the 4.0 release (the current version is 7.1), the *mail()* function is yet to be fixed after 15 years. However, it is specified in the PHP documentation [15] that the *mail()* function does not protect against this vulnerability. A working code sample of the vulnerability, written in PHP 5.6 (latest well-supported version), is shown in Listing 1.

2.3.2 Python

A bug was filed about the vulnerability in Python's implementation of the *email.header* library and its header parsing functions allowing newlines in early 2009, which was followed up with a partial patch in early 2011.

Unfortunately, the bug fix was only for the *email.header* package, and thus is still prevalent in other frequently used packages such as *email.parser*, where both the classic *Parser()* and the newer *FeedParser()* exhibit the vulnerability even in the latest versions - 2.7.11 and 3.5. The bug fix was also not backported to older versions of Python. There is no mention of the vulnerability in the Python documentation for either library. A working code sample of the vulnerability, written in Python 2.7.11, is shown in Listing 3.

```
from email.parser import Parser
import cgi
form = cgi.FieldStorage()
to = form["email"] # input() exhibits the
                    same behavior
msg = """To: """ + to + """\n
From: <user@example.com>\n
Subject: Test message\n\n
Body would go here\n"""
f = FeedParser() # Parser.parsestr() also
# contains the same vulnerability
f.feed(msg)
headers = FeedParser.close(f)
# attack string => 'sai@sai.com\nBCC:
# spc@spc.com'
# for form["email"]
# both to:sai@sai.com AND bcc:spc@spc.com
# are added to the headers
print 'To: %s' % headers['to']
print 'BCC: %s' % headers['bcc']
```

Listing 3: Python program with e-mail header injection vulnerability.

2.3.3 Java

Java has a bug report about E-Mail Header Injection filed against its JavaMail API. A detailed write-up by Alexandre Herzog [6] is complete with a proof of concept program that exploits the API to inject headers.

2.3.4 Ruby

From our preliminary testing, Ruby's built-in Net::SMTP library has this vulnerability. This is not documented on the library's homepage. A working code sample of the vulnerability, written in Ruby 2.0.0 (the latest stable version at the time of writing), is shown in Listing 4.

```
1 require 'sinatra'
2 require 'net/smtp'
3
4 get '/hello' do
```

Server Side Language	% of Usage
PHP	81.9
ASP.NET	15.8
Java	3.1
Ruby	0.6
Perl	0.5
JavaScript	0.2
Python	0.2

Table 2: Language usage statistics compiled from w3techs [22].

```

5 email = params[:email]
6
7 message = ""
8 From: Sai <schand31@asu.edu>
9 Subject: SMTP e-mail test
10 To: #{email}
11
12 This is a test e-mail message.
13 ""
14 # construct a post request with email set
15 # to attack_string
16 # attack_string => sai@sai.com%0abcc:
17 # spc@spc.com%0aSubject: Hello
18 Net::SMTP.start('localhost', 1025) do |smtp|
19 |
20 | smtp.send_message message, 'schand31@asu.edu',
21 | 'to@todomain.com'
22 end
23 # Headers get added, and Subject field
24 # changes to what we set.
25 end

```

Listing 4: Ruby program with e-mail header injection vulnerability.

2.4 Potential Impact

The impact of the vulnerability can be pretty far-reaching. Table 2 shows the current server-side language usage statistics on the Web [22].

PHP, Java, Python, and Ruby (combined) account for over 85%¹ of the websites measured. The vulnerability can be exploited to do potentially any of the following:

- **Phishing and Spoofing Attacks**
Phishing [14] (a variation of spoofing [28]) refers to an attack where the recipient of an e-mail is made to believe that the e-mail is a legitimate one. The e-mail usually redirects them to a malicious website, which then steals their credentials.
E-Mail Header Injection gives attackers the ability to inject arbitrary headers into an e-mail sent by a website and control the output of the e-mail. This adds credibility to the generated e-mail, as it is sent right from the websites and people are more ready to trust e-mail that is received from the website directly and can thus result in more successful phishing attacks.
- **Spam Networks**
Spam networks can use E-Mail Header Injection vulnerabilities on the ability to send a large amount of

¹A website may use more than one server-side programming language

e-mail from servers that are trusted. By adding additional “cc” or “bcc” headers to the generated e-mail, attackers can easily achieve this effect.

Due to the e-mails being from trusted domains, recipient e-mail clients might not flag them as spam. If they do flag them as spam, then that can lead to the website being blacklisted as a spam generator.

- **Information Extraction of legitimate users**
E-Mails can contain sensitive data that is meant to be accessed only by the user. Due to E-Mail Header Injection, an attacker can easily add a “bcc” header, and send the e-mail to himself, thereby extracting important information. User privacy can thus be compromised, and loss of private information can by itself lead to other attacks.
- **Denial of service by attacking the underlying mail server**
Denial of service attacks (DoS), can also be aided by E-Mail Header Injection. The ability to send hundreds of thousands of e-mails by just injecting one header field can result in overloading the mail server, and cause crashes and/or instability.

It is evident that E-Mail Header Injection is a critical vulnerability that web applications must address.

3. SYSTEM DESIGN

This chapter discusses the System design, and explains the architecture and the components of the System in detail. It then proceeds to enumerate the issues faced and the assumptions made during the building of the system.

3.1 Our Approach to measure prevalence of E-Mail Header Injection

We took a black-box approach to measure the prevalence of E-Mail Header Injection vulnerability on the World Wide Web. Black-box testing [23] is a way to examine the functionality of an application without looking at its source code.

As we did not have the source code for each of these websites, black-box testing was the ideal approach for this project. Black-box testing allows our system to detect E-Mail Header Injection vulnerabilities in any server-side language (not simply those we identified in Section 2.3). A high level logical overview of our system is presented in Figure 1. The components shown in Figure 1 can be more broadly categorized into different modules, as discussed in the following section.

3.2 System Architecture

The black-box testing system can be divided broadly into two modules; Data Gathering and Payload Injection.

1. Data Gathering

The Data Gathering module (shown in Figure 2) is primarily responsible for the following activities:

- Interface with the Crawler (Section 3.3.1) and receive the URLs.
- Parse the HTML for the corresponding URL and store the relevant form data (Section 3.3.2).
- Check for the presence of forms that allow the user to send/receive e-mail, and store references to these forms (Section 3.3.3).

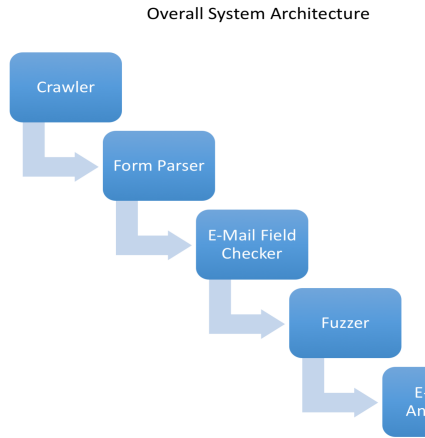


Figure 1: Overall system architecture - logical overview.

2. Payload Injection

The Payload Injection module (shown in Figure 3) is primarily responsible for the following activities:

- Retrieve the forms that allow users of a website to send/receive e-mail and reconstruct these forms (Section 3.3.4).
- Inject these forms with benign data (non-malicious payloads) and generate an HTTP request to the corresponding URL (Section 3.3.5).
- Analyze the e-mails, extracting the header fields and checking for the presence of the injected payloads (Section 3.3.6).
- Inject the forms that sent us e-mails with malicious payloads, and generate an HTTP request to the corresponding URL to check if E-Mail Header Injection vulnerability exists in that form (Section 3.3.5).

The functionality of each component is discussed further in the ‘Components’ section (Section 3.3). The Payload Injection pipeline is not a linear, but cyclic process, as we inject different payloads and analyze the received e-mails.

3.3 System Components

The Data Gathering module and Payload Injection module are made up of a number of smaller components. This section describes in detail the functionality of each of the components.

3.3.1 Crawler

We used an open-source Apache Nutch based Crawler. The Crawler provides us with a continuous feed of URLs and the HTML contained in those pages. This feed is sent to our Form Parser over a Celery Queue.

3.3.2 Form Parser

The actual pipeline begins at the Form Parser. This module is responsible for parsing the HTML and retrieving data about the forms on the page, including the following:

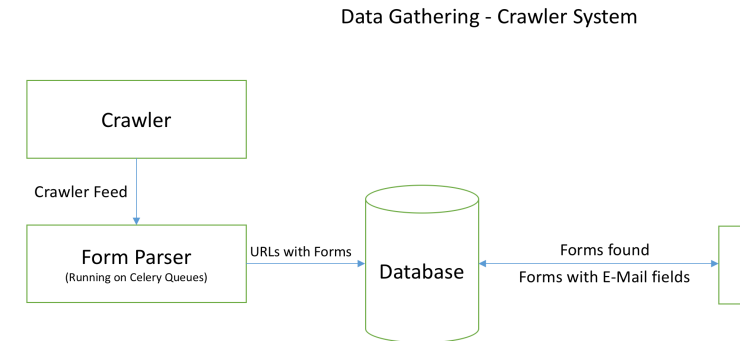


Figure 2: System architecture - crawler & form parser.

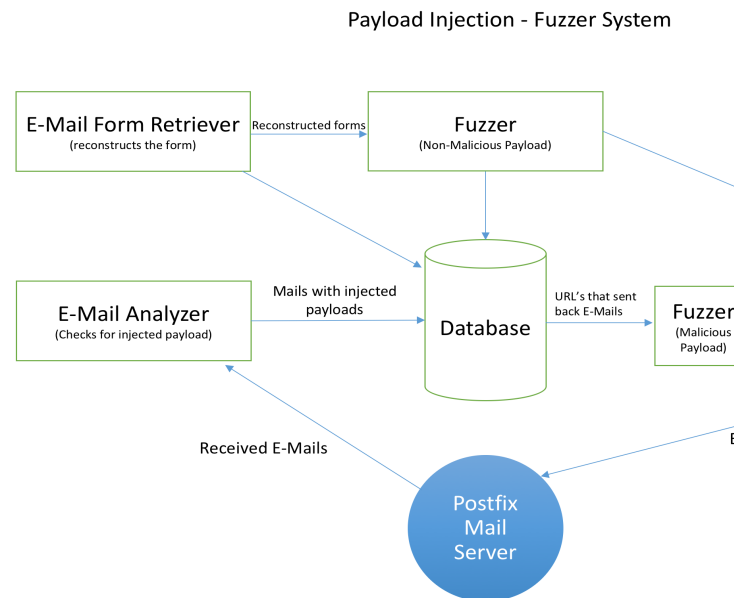


Figure 3: System architecture - fuzzer & e-mail analyzer.

- Form attributes, such as method and action. These dictate where we send the HTTP request and what kind of request it is (GET or POST).
- Data about the input fields, such as their attributes, names, and default values. The default values are essential for fields like `<input type="hidden">` as these fields are usually used to check for the submission of forms by bots.
- Presence of the `<base>` element in the HTML, as this affects the final URL to which the form is to be submitted.
- Headers associated with the page, such as *referrer*. Once again, these were required to avoid the website from ignoring our system as a bot.

The Form Parser stores all this data in our database, so as to allow us to reconstruct the forms later for fuzzing, as required.

3.3.3 E-Mail Field Checker

The E-Mail Field Checker script is the final stage in the Data Gathering pipeline. It receives the output of the previous stage—form data from the queue—and checks for the presence of e-mail fields in those forms. If any e-mail fields are found, it stores references to these forms in a separate table. This separates the forms that are potentially vulnerable from the forms that are not.

The E-Mail Field Checker searches for the words ‘e-mail’, ‘mail’ or ‘email’ within the form, instead of an explicit e-mail field (e.g., `<input type="email">`). This is by design, taking into account a very common design pattern used by web developers, where they may have a text field with an “id” or “name” set to ‘email’, instead of an actual e-mail field, for purposes of backward compatibility with older browsers.

Compared to searching for explicit e-mail fields, by searching for the presence of the words ‘e-mail’, ‘mail’ or ‘email’ in the form, we are assured very few false negatives. This is because our system is bound to find e-mail fields with their “type”, “name”, or “id” set to one of these words. The system is also substantially faster as we do not have to parse the individual form fields at this point in the pipeline. However, despite the advantages, this might also lead to a false positive rate. We discuss this possibility in detail in Section 3.4 - Design Issues.

The output of this stage is stored in the database for persistence and acts as the input to the ‘Payload Injection’ pipeline.

3.3.4 E-Mail Form Retriever

The E-Mail Form Retriever is the first stage in the Payload Injection Pipeline. It has the following important functions:

- Retrieve the newly inserted forms in the ‘email_forms’ table, checking to ensure no duplication occurs before the fuzzing stage.
- Reconstruct each form, using the data stored in the ‘form’ table, complete with input fields and their values.
- Construct the URL for the ‘action’ attribute of the form so that we can send the HTTP request to the correct URL.

3.3.5 Fuzzer

The Fuzzer is the heart of the system and is the only component that interacts directly with the external websites. The Fuzzer is split into smaller modules, each of which is responsible for a particular type of fuzzing. We inject payloads in two different stages, to improve the efficiency, and reduce the total number of HTTP requests we generate. This is because making HTTP requests is an expensive process [12], and can be a cause of bottlenecks in a Crawler-Fuzzer system [18]. The two different types of payloads we use for fuzzing are:

Non-Malicious Payload.

The regular or non-malicious payload is a straight forward E-Mail address of the format – ‘reguser(xxxx)@example.com’, where ‘xxxx’ is replaced by our internal ‘form_id’, to create a one-to-one mapping of the payloads to the forms, and ‘example.com’ is replaced by the required domain. In our case, this domain was ‘wackopicko.com’. This non-malicious payload allows us to check whether we can inject data into a form and whether we can overcome the ‘anti-bot’ measures on the given website, without attempting to fuzz the website.

Malicious Payload.

In the malicious payload scenario, we inject the fields with the “bcc” (blind carbon copy) element. If the vulnerability is present, this will cause the server to send a copy of the e-mail to the e-mail address we added as part of the ‘bcc’ field.

We consider a special case: the addition of a “x-check:in” header field to the payloads. This is due to Python’s exhibited behavior when attaching headers. Instead of overwriting a header if it is already present, it ignores duplicate headers. So, in case the “bcc” field is already present as part of the headers, our injected “bcc” header would be ignored. To overcome this, we need to inject a new header that is not likely to be generated by the web application. Hence, we inject our own “x-check:in” header to ensure we can get results if the injection was successful.

The malicious payloads consist of 4 different payloads. Each of these payloads is crafted for a particular use case. The four payloads are:

1. `nuser(xxxx)@wackopicko.com\nbcc:maluser(xxxx)
@wackopicko.com`
- This is the most minimal payload, it injects a ‘newline’ character followed by the ‘bcc’ field.
2. `nuser(xxxx)@wackopicko.com\r\nbcc:maluser(xxxx)
@wackopicko.com`
- This payload is added for purposes of cross-platform fuzzing: ‘\r\n’ is the ‘Carriage Return - New Line (CRLF)’ used on Windows systems.
3. `nuser(xxxx)@wackopicko.com\nbcc:maluser(xxxx)
3@wackopicko.com\nx-check:in`
- As discussed above, the addition of the “x-check:in” header is to inject Python based websites.
4. `nuser(xxxx)@wackopicko.com\r\nbcc:maluser(xxxx)
4@wackopicko.com\r\nx-check:in`
- Same as the previous payload, but containing the additional ‘\r’ for Windows compatibility.

The ‘xxxx’ in all of the payloads is replaced by our internal ‘form_id’, so as to create a one-to-one mapping of the payloads to the forms. The coverage provided by each payload is shown in Table 3.

Payload	Languages covered	Platforms covered
1	PHP, Java, Ruby, etc.	Unix
2	PHP, Java, Ruby, etc.	Windows
3	Python	Unix
4	Python	Windows

Table 3: Payload coverage, each payload covers a different platform/language.

Along with the payload, the Fuzzer also injects data into the other fields of the form. This data must pass validation constraints on the individual input fields e.g., for a name field, numbers might not be allowed. It is essential that the data we inject into the input fields adhere to the constraints. Our Fuzzer does this by making use of a ‘Data Dictionary’ which has predefined ‘keys’ and ‘values’ for standard input fields such as **name**, **date**, **username**, **password**, **text**, and **submit**. The default values for these are generated on-the-fly for each form, based on generally followed guidelines for such fields. For example, password fields should consist of at least one uppercase letter, one lowercase letter, and a special character.

Once the data (including the payload) for the form is ready, the Fuzzer constructs the appropriate HTTP request (GET or POST) and sends the HTTP request to the URL that was generated by the E-Mail Form Retriever (Section 3.3.4).

3.3.6 E-Mail Analyzer

The E-Mail Analyzer checks for the presence of injected data in the received e-mails. This module works on the e-mails received and stored by our Postfix server, and depending on the user who received the e-mail, it performs different functions.

Analyzing regular e-mail.

‘Regular e-mail’ refers to the e-mails received by the `reguser(XXXX)@wackopicko.com` — where ‘XXXX’ is our internal ‘form_id’ — that were sent due to injecting the ‘regular or non-malicious’ payload (discussed in Section 3.3.5). The objective of the analysis on this e-mail is identify if the input fields that we injected with data appear on the resulting e-mail, and if so, which fields appear where.

To find this, we read through each received e-mail, and check whether *any* of the fields we injected with data appear as part of either the headers or the body of the e-mail. If they do, we add them to the list of fields that can potentially result in an E-Mail Header Injection for the given e-mail. We then pass on this information back to the Fuzzer pipeline, along with the ‘form_id’, so that the Fuzzer can now inject the malicious payloads into the same form.

Analyzing e-mail with payloads.

The ‘e-mails with payloads’ refer to e-mails received by either the `nuser(XXXX)@wackopicko.com` or `maluser(XXXX)@wackopicko.com` accounts. These e-mails were received due to injecting the malicious payloads that were discussed in Section 3.3.5.

Analysis of these e-mails is considerably simpler than that of the regular e-mails. This is due to the fact that this involves lesser processing of the contents of the e-mail compared to the previous section.

Detecting injected bcc headers.

As discussed in the payloads section (3.3.5), the payloads were crafted in such a way that the e-mails received by ‘maluser’ account directly indicate the presence of the injected ‘bcc’ field. Thus, we simply parse the E-Mails and store them in the Database.

Detecting injected x-check headers.

E-Mails not received by the ‘maluser’ account but by the ‘nuser’ account constitute a special category of e-mails. These e-mails could have been generated due to two reasons:

1. The websites performed some sanitization routines and stripped out the “bcc” part of the payload, thereby sending e-mails only to the ‘nuser’ account. These e-mails then act as proof that the vulnerability was not found on the given website.
2. A more conducive scenario is when the “bcc” header was ignored for some reason, e.g. Python’s default behavior when it encounters duplicate headers. In this case, we check whether the e-mail contains the custom header “x-check”. If it does, then this is a successful exploit of the vulnerability, and we store it in the database.

3.3.7 Database

We collect and store as much data as possible at each stage of the pipeline. This is due to the two following reasons:

1. The data is used to validate our findings.
2. The data collected can be used for other research projects in this area.

Each table in our database is listed in Table 4 along with the data it is designed to hold. A schema of the database is shown in Figure 4.

3.4 Design Issues

This section will describe the issues we faced with the design decisions we made, and how we did our best to mitigate them, and their effect on the system.

- False Positive rate for the E-Mail Field Checker

As discussed in Section 3.3.3, we only search for the words ‘email’, ‘mail’ or ‘e-mail’ (case insensitive) inside the forms to detect the presence of e-mail fields, instead of searching for an `<input type = email>`. This might result in a false positive in certain forms, like the one shown in Listing 5.

```

1  <form method="post">
2  E-Mail us if you have any
   questions!!
3  <input type="text" name="query"
   <br>
4  <input type="submit" value="
   Search">
5  </form>
```

S.No	Table Name	Purpose
1	form	To hold data about all the forms that we receive from the Form Parser.
2	email_forms	Holds the output of the E-Mail Field Checker, i.e. references to the ID's of the forms that contain E-Mail fields.
3	params	Holds the actual input fields of the forms, including their default values.
4	fuzzed_forms	Holds the data of the forms that were injected, including the payload used to inject and the URL to which the HTTP Request was delivered.
5	received_emails	Contains data about the E-Mails received for the malicious payload. This contains the end result of the proposed injection pipeline.
6	successful_attack_emails	Contains data about the E-Mails received for the malicious payload. This contains the end result of the proposed injection pipeline.
7	requests	Contains data about the requests generated for each URL.
8	blacklisted_urls	Used for skipping certain websites that are blacklisted by the Crawler-Fuzzer.

Table 4: The different tables in our database.

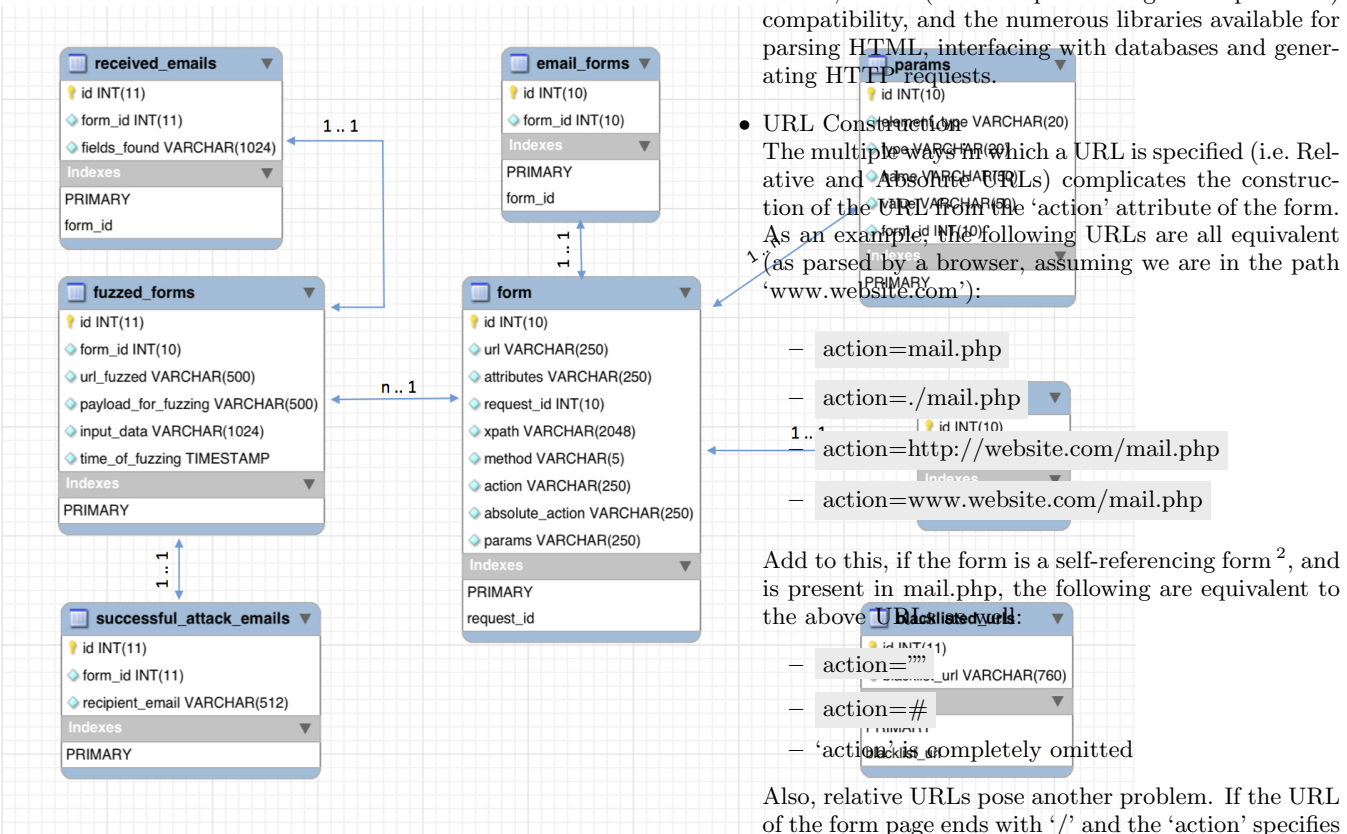


Figure 4: Database schema.

Listing 5: E-Mail field checker – false positives, the system incorrectly classifies this as an e-mail form.

The word ‘E-Mail’ on Line 2 will result in our system classifying this form as a potential e-mail form, while it clearly is not. However, as we will see, this is not really a significant issue, as despite being added to the ‘email_forms’ table, this form will never be injected in the ‘fuzzer’ due to the absence of the appropriate input field in the form. We chose to go with this design, as it allows us to detect almost every form that provides the capability to send or receive e-mail.

Parallelism for the system. Every component in the pipeline benefits hugely from parallel processing of the data. However, Python’s GIL (Global Interpreter Lock) for each URL. To overcome this, we used a Celery task queue (discussed in Section 4.4), which allowed a level of parallelism that Python does not provide by default. Even though this makes the system faster than a single-threaded approach, it still leaves room for improvement in terms of performance. Despite the speed drop that results from lack of full parallelism, we chose to go with Python, for the raw power it provides, its text processing capabilities, PCRE (Perl Compatible Regular Expressions) compatibility, and the numerous libraries available for parsing HTML, interfacing with databases and generating HTTP requests.

• URL Construction. The multiple ways in which a URL is specified (i.e. Relative and Absolute URLs) complicates the construction of the URL from the ‘action’ attribute of the form. As an example, the following URLs are all equivalent (as parsed by a browser, assuming we are in the path ‘www.website.com’):

- action=mail.php
- action=./mail.php
- action=http://website.com/mail.php
- action=www.website.com/mail.php

Add to this, if the form is a self-referencing form², and is present in mail.php, the following are equivalent to the above URLs:

- action=
- action=#
- ‘action’ is completely omitted

Also, relative URLs pose another problem. If the URL of the form page ends with ‘/’ and the ‘action’ specifies a path starting with ‘/’ (illustrated in Listing 6), we

²A self-referencing form is one which submits the form data to itself. It includes logic to both display the form and process it. It is a *very* common feature in PHP-based scripts.

would need to strip one of the two slashes. This increases the overall complexity of our URL generator, as we have to account for all these possibilities.

```
1 Current URL = www.website.com/  
2 <form action=/mail.php>
```

Listing 6: URL construction, the resulting url needs to be `www.website.com/mail.php` and not `www.website.com//mail.php`

As using a browser engine to reconstruct these URLs and connecting it to the fuzzer pipeline would have added unnecessary bulk to the project, we chose to go with a best-effort approach to this problem, where our system covers all these possibilities with a lightweight URL Generator, however, we cannot know for certain whether this works for other unforeseen ways of specifying a URL.

- **Black-box Testing**

The approach that we have selected — Black-box testing — is highly beneficial as explained in Section 3.1. However, it also has a drawback in that we cannot verify whether the reported vulnerability exists in the source code or is a feature of the website (e.g., the website allows users to send bulk e-mail, adding as many “cc” or “bcc” headers). We have to manually e-mail the developers to get this feedback.

- **Mapping responses to requests**

As we are generating multiple payloads for each form, and the received e-mail may not contain the name of the domain from which we received the e-mail, it is difficult to map the response e-mails to the right requests. We instead use the ‘form_id’ as part of the payload to map responses to requests accurately.

- **Bot Blockers**

Because our system is fully automated, it is also susceptible to being stopped by ‘bot-blockers’ i.e. mechanisms built-in to a website to prevent automated crawls or form submissions. Measures like CAPTCHA (Completely Automated Public Turing test to tell Computers and Humans Apart) and hidden form fields are often used to detect bots [16], [21].

We have made sure that we do not affect hidden fields in the form, however, we do not have an anti-CAPTCHA functionality built into our system, and thus our system will not test such websites.

- **Handling Malformed HTML**

The parser that we use for HTML parsing — Beautiful Soup — does not try to parse malformed HTML, and throws an exception on encountering malformed content. Thus, we have designed the system to exit gracefully on such occasions. A side-effect of this is that our system is unable to parse websites with bad markup³.

³We do not have any data about whether bad markup indicates an overall lower quality of the website, and thus cannot comment on whether such websites are more likely to have vulnerabilities, although the author strongly suspects that that might be the case.

- **Crawling WordPress and other CMS-based websites**

In contrast to bot blockers that try to prevent the automated systems from attacking them, WordPress and other CMS based websites use a blacklisting approach to prevent bot attacks. Unfortunately, because we generate multiple requests to each website, this results in our IPs getting blacklisted. To overcome this, we did two things:

1. Used an IP range of 60 different IP addresses.
2. Used a blacklist of our own to prevent our Fuzzer from fuzzing websites that are known to blacklist automated crawlers.

3.5 Assumptions

We made certain assumptions while building the system. This section describes the assumptions and explores to what extent these hold true:

1. **Crawler is not blocked by firewalls**

This is a requisite for our system to work. If the Crawler is blocked for any reason, we do not get the data feed for our system, and without this input, it is almost impossible to set our system up.

2. **The Crawler feed is an ideal representation of the World Wide Web**

This is a reasonable expectation, albeit an unrealistic one.

It is unrealistic because Crawlers work on the concept of proximity. They detect for the presence of In-Links and Out-Links from a particular URL, and hence the returned URLs are usually related to each other (at least the ones that are returned adjacent to each other).

However, this assumption is reasonable due to the ‘Law of averages’ [24], the ‘Law of big numbers’ [25], and the concept of ‘Regression to the mean’ [27]. Simply stated, a crawl of this large magnitude should give us a very distributed sample of the overall Web, eventually converging to the average of all websites in existence.

3. **Injection of “bcc” indicates the existence of E-Mail Header Injection Vulnerability**

We assume that the ability to inject a “bcc” header field is proof that the E-Mail Header Injection vulnerability exists in the application. We do not inject any additional payloads that can modify the subject, message body, etc. as this analysis is designed to be as benign as possible. We believe that this is a reasonable assumption, as altering e-mail headers is a goal of exploiting E-Mail Header Injection vulnerability.

That concludes our discussion about the design of the system. To recap, we discussed our approach, the system architecture and how the components fit into our architecture. We also discussed the issues faced, and the assumptions that we made while building the system.

4. EVALUATION

This chapter describes the experimental setup for our project including the servers used, the software and the platforms involved, the languages used, and the task queue system that was used for parallelism. We follow this up with our evaluation of the system, with a test suite, and proof of concept examples.

4.1 System Configuration

We used two systems for the project, and their configurations are as follows:

- **Dell PowerEdge T110 II Server**

CPU: Intel(R) Xeon(R) CPU E3-1220 V2 @ 3.10GHz
Cache size : 8192 KB
No. of Cores : 4
Total Memory (RAM) : 16 GB
Disk Space : 2 TB

- **MacBook Pro**

CPU: Intel Core i7 @ 2.8 GHz
Cache size : 6144 KB
No. of Cores : 4
Total Memory (RAM) : 16 GB
Disk Space : 500 GB

4.2 Platforms and Software

We enumerate the platforms and the software used for our project in Table 5.

Operating system	Ubuntu 14.04
Server	Apache - 2.4.17
Database	MariaDB - 10.1.9
Mail Server	Postfix - 2.11.0
Other software used	Mailcatcher, PostMan, HTTPRequester, RabbitMQ

Table 5: Platforms and software used for our project.

4.3 Languages Used

We used Python 2 to build the system. The following factors influenced our choice of language: text processing capabilities, PCRE (Perl Compatible Regular Expressions) compatibility, and the numerous libraries for HTML Parsing, HTTP request generation, mail processing etc. We made use of the following major libraries (shown in Table 6) for our system.

Library	Functionality
Requests	HTTP Request Generation
Beautiful Soup	HTML Parsing
Mailbox	Mail Processing
Celery	Task Queues

Table 6: Libraries that we used and their functions.

Despite the many benefits that Python 2 provides, we had certain issues with the language — discussed in Section 3.4 — such as Python’s GIL (Global Interpreter Lock) which does not allow the running of multiple native threads concurrently. The following section (Section 4.4) describes in detail the task queue system (Celery) that we used to overcome this limitation of Python.

4.4 Celery Queues

We used a Celery task queue running on RabbitMQ to overcome the GIL. According to Celery Project Homepage [3]:

“Celery is an asynchronous task queue/job queue based on distributed message passing.”

Simply put, Celery allows us to process multiple tasks in parallel by making use of what is known as a task queue. Celery instantiates multiple workers that listen to these queues and processes each task individually. This simulates pseudo-parallel processing to a certain degree, by allowing us to run multiple instances of the same program. It does this by using a message broker called RabbitMQ. According to RabbitMQ’s Wikipedia page [26],

“RabbitMQ is an open source message broker software that implements the Advanced Message Queuing Protocol (AMQP)”

RabbitMQ facilitates the storage and transport of messages on queuing systems. It is also cross-platform and open source, providing us with clients and servers for many different languages, thereby being the ideal fit for Celery. Thus, by using Celery and RabbitMQ together, we were able to achieve a certain degree of parallelism that would not have been possible with traditional Python.

4.5 Test Suite

The test plan for our system includes a set of unit tests for each module in the pipeline. Further, we have unit tests for every individual function in the modules. The functions are tested separately, using mocks and stubs, so as to ensure isolated testing. This section outlines the test plan in the following manner. We list the modules that are tested, and then describe what each unit test tests for.

- **Form Parser**

- `test_url_exception` - Tests whether the system handles incorrect or malformed URLs properly and terminates cleanly.
- `test_db_connection` - Tests whether the Database Connection is set up and queries can be executed.
- `test_form_parser` - Tests for the proper parsing of HTML, and if the system exits cleanly in case parsing is not possible.

- **E-Mail Field Checker**

- `test_check_for_email` - Tests whether the system finds E-Mail fields in the form when the words ‘e-mail’ or ‘email’ are present in the form (case insensitive).
- `test_check_for_no_email` - Tests whether the system finds no E-Mail fields when the words ‘e-mail’ or ‘email’ are *not* present in the form (case insensitive).

- **E-Mail Form Retriever**

- `test_reconstruct_form` - Tests for the proper reconstruction of the form stored in the Database.

- `test_construct_url` - Tests whether the URL for submission was constructed properly, includes checks for relative URLs, absolute URLs, and presence of 'base' tags.
 - `test_email_form_retriever_already_fuzzed` - Tests for duplicate fuzz requests, and whether the system rejects these requests.
 - `test_email_form_retriever_calls_fuzzer_for_new_fuzz` - Tests whether the E-Mail Form Retriever calls the Fuzzer module with the proper data when it gets a new fuzz request.
- Fuzzer
 - `test_send_get_request` - Tests for the proper handling of GET requests.
 - `test_send_post_request` - Tests for the proper handling of POST requests.
 - `test_correct_fuzzer_data` - Tests whether the payload generated for the given form data is correct and consistent. Also tests whether the payload was part of the resulting HTTP request.
 - `test_incorrect_fuzzer_data` - Tests for incorrect form data, and ensures that a payload does not end up in the wrong input field in the resulting HTTP request.
 - E-Mail Analyzer
 - `test_load_mail` - Tests whether the E-Mails are loaded and parsed correctly by the E-Mail Analyzer.
 - `test_parse_headers` - Tests for the proper parsing of headers present in the E-Mail.
 - `test_analyze_regular_mail` - Tests whether the E-Mail Analyzer parses the regular E-Mail properly and extracts the injected input fields that are present in the E-Mail.
 - `test_analyze_malicious_mail` - Tests whether the E-Mail Analyzer parses the E-Mails received due to the malicious payloads properly, is able to extract the 'bcc' headers, and is able to link them to the proper fuzzing request and payload.
 - `test_analyze_x_check_header` - Tests whether the 'x-check' header is read by the E-Mail Analyzer.

The unit tests were written using Python's built-in 'Unittest' module, mocking was done using the built-in 'MagicMock' module. The tests allow us to be reasonably certain that our system works as expected.

4.6 Proof of Concept Attacks

The previous section (Section 4.5) discussed in detail how we verified that our system functions according to our expectations. This section describes how we validated our expectations. In order to do this, we constructed three sets of web applications in PHP, Python, and Ruby. Each of these applications was a simple web app that accepted user input to construct and send an E-Mail.

The front-end for each of the three applications is shown in Listing 7. The server-side code for the three languages are shown in Listings 1, 3, and 4.

We tested for the headers being injected in real-time by running an instance of MailCatcher, set to listen on all SMTP messages. A sample screenshot of a fuzzed request for the Ruby backend (generated in PostMan) is shown in Figure 5. The e-mail sent due to injecting this payload (as captured by MailCatcher) is shown in Figure 6. It can be seen that the headers have been added to the resulting e-mail, and we have successfully managed to overwrite the "Subject" field with our message, 'hello'.

A similar injection example for PHP is shown in Figure 7 and the corresponding e-mail caught by MailCatcher is shown in Figure 8. The astute reader might have noticed that in the given examples we have used '%0a' to separate the headers, while in Section 3.3.5, we had used '\n'. This is due to URL encoding [1], wherein special characters in the URL are 'encoded' or 'escaped' with their ASCII equivalent. The reason why we do not have to do this with the payloads our system injects is due to the fact that the Python Requests library that we use to generate the HTTP requests automatically does this encoding for us.

```

1 <!doctype html>
2 <html lang="en">
3 <head>
4 <meta charset="utf-8">
5 <meta name="author" content="Sai Pc">
6 <title>Mock Email</title>
7 </head>
8 <body>
9 <form action="{Replace with path to back-
  end}" method="post">
10 <input type="text" placeholder="Email" name
   ="email" id="e-mail"><br>
11 <textarea name="msg" rows="20" cols="120"><
   /textarea>
12 <input type="submit" value="Email Me!">
13 </form>
14 </body>
15 </html>

```

Listing 7: HTML page for showcasing e-mail header injection, a simple front-end for our examples.

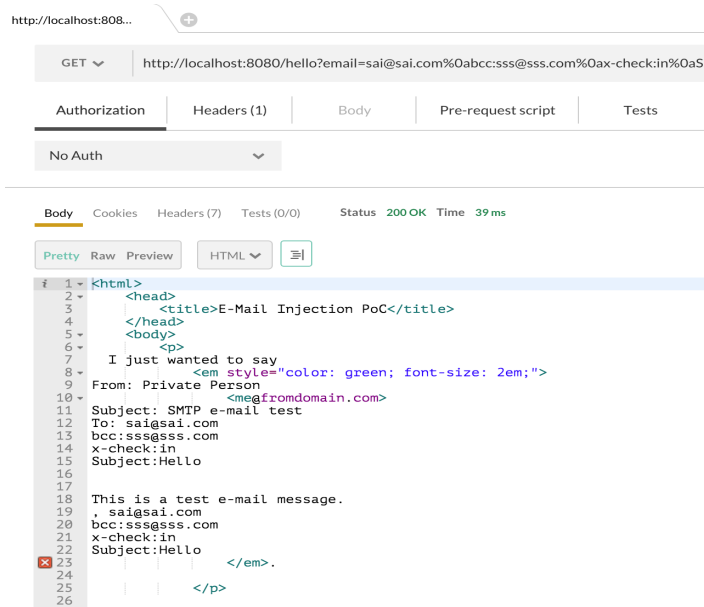


Figure 5: Fuzzing a request for the Ruby backend, the payload can be seen inside the address bar.

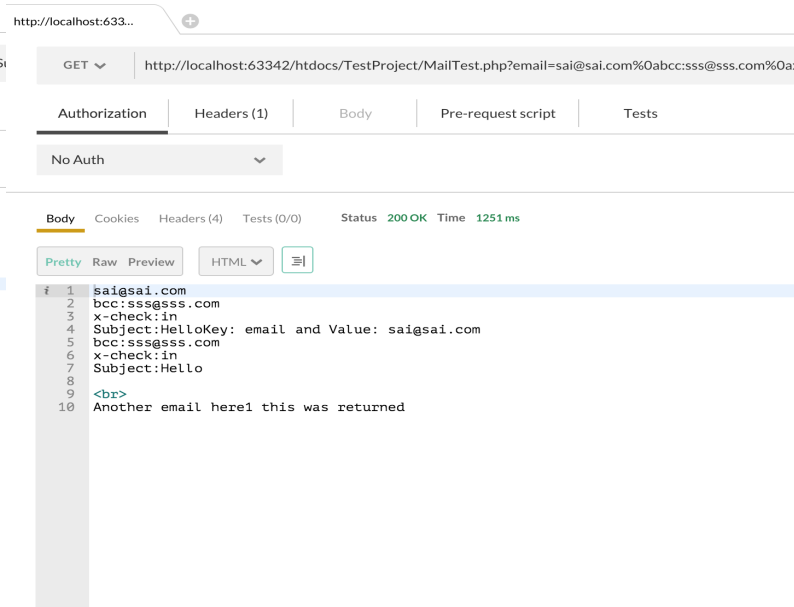


Figure 7: Fuzzing a request for the PHP backend, the payload can be seen inside the address bar on top.

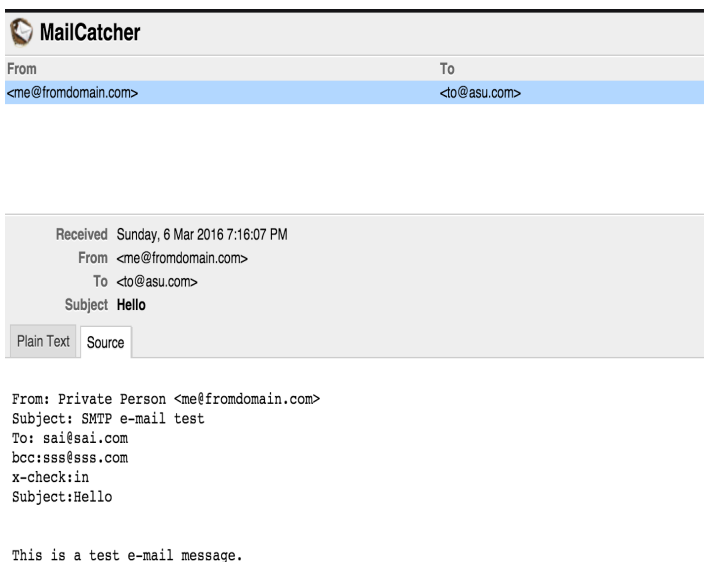


Figure 6: E-Mail header injection proof of concept - Ruby, we can see that multiple headers (bcc, x-check, subject) have been inserted into the resulting e-mail.

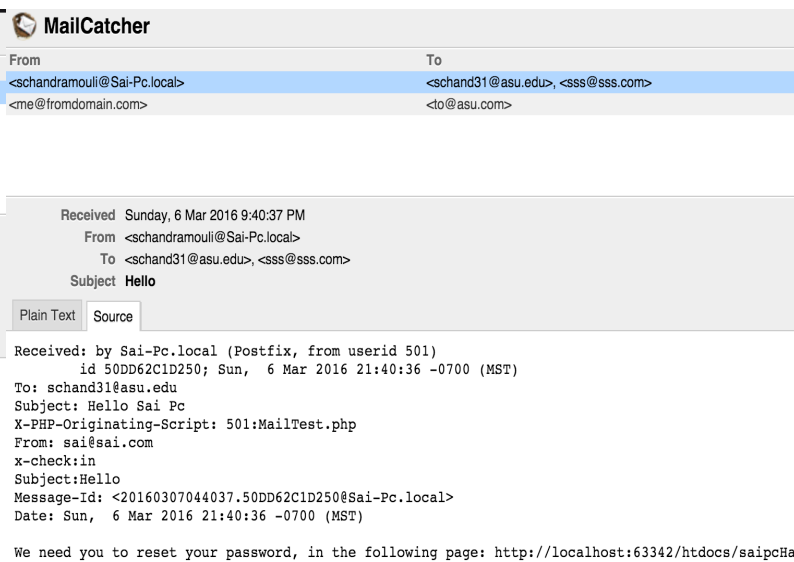


Figure 8: E-Mail header injection proof of concept - PHP, we can see that multiple headers (bcc, x-check, subject) have been inserted into the resulting e-mail.

5. DATA ANALYSIS AND RESULTS

This chapter serves to present our findings: the data that we gathered from our crawl, the data generated due to the fuzzing attempts and our analysis on this data.

5.1 Collected Data

From our extensive crawl of the web, we were able to gather the data shown in Table 7. The following paragraphs describe in detail what each kind of data represents, and what they signify.

S.No	Type of Data	Quantity
1	URLs Crawled	21,675,680
2	Total Forms found	6,794,917
3	Forms with E-Mail Fields	1,132,157

Table 7: The data collected for our project.

URL's crawled.

This represents the total number of unique URLs that we crawled on the World Wide Web. It is to be noted that this quantity refers to unique URL's, and not websites.

Total Forms found.

This represents the total number of forms that were found on the URL's crawled. We found a total of 6,794,917 forms from 1,019,921 unique domains.

Forms with E-Mail Fields.

This represents the total number of forms that were found to contain e-mail fields. We found 1,132,157 such forms from 197,570 unique domains.

5.2 Fuzzed data

We performed our fuzzing attempts on the gathered data with both the regular payload and malicious payload. Table 8 shows the quantity of e-mails that we received for each payload. We explain in detail what each piece of data shown in the table represents, in the following sections.

S.No	Type of fuzzing	Forms fuzzed	E-Mails received
1	Regular payload	934,016	52,724
2	Malicious payload	46,156	496

Table 8: The data that we fuzzed and the e-mails that we received.

E-Mail received from forms.

The e-mails that we received can be broadly categorized into two categories:

1. E-Mails due to regular payload
This represents the total number of websites that sent e-mails to us. This indicates that we were able to successfully submit the forms on these sites.
2. E-Mails due to malicious payload
Once we receive an e-mail from a website due to the regular payload, we go back and fuzz those forms with more malicious payloads. This field, in essence, represents the total number of unique URLs that contain E-Mail Header Injection vulnerability.

5.3 Our Analysis of the received e-mail Data

During our analysis of the received e-mails, we found that the e-mails that we received belonged to one of three broad categories:

1. E-Mails with the “bcc” header successfully injected
This form of injection was our initial objective and we found 265 such e-mails in our received e-mails. This indicates that the websites that sent out these e-mails are vulnerable to e-mail header injection, where we could inject and manipulate any header.
2. E-Mails with the “to” header successfully injected
We discovered an unintended vulnerability which we would like to christen ‘TO header injection’. These injections reflect the ability to inject any number of e-mail addresses into the “to” field while being unable to inject any other header into the e-mails. We attribute this behavior to inconsistent sanitization by the application. The vulnerability is further aided by the leniency shown by mail servers, wherein they parsed malformed e-mail addresses and delivered it to the right mail server, and on the receiving end, the mail was delivered to the right mailbox.

While not allowing us complete control over the e-mails sent, ‘TO header injection’ makes it possible to append any number of e-mail addresses, thereby enabling us to leak information, and/or perform DoS (Denial of Service) attacks.
3. E-Mails with the “x-check” header successfully injected
The third category of e-mails received were e-mails with the “x-check” header injected. As discussed in Section 3.3.6, these let us differentiate between unsuccessful attempts and successful attempts by injecting the additional header, allowing us to check whether headers other than the “bcc” header can be injected into the generated e-mail.

We list each category and the number of e-mails received by the category in Table 9.

S.No	Type of Injection	N
1	E-Mail Header Injections with ‘bcc’ header	
2	E-Mail Header Injections with ‘x-check’ header	
3	‘To header’ injections alone	
4	E-Mail Header Injections with ‘bcc’ and ‘x-check’ headers	
5	Both ‘To header’ injections and x-check headers	
6	‘x-check’ headers found in ‘nuser’ e-mails	
7	Unique ‘x-check’ headers found in ‘nuser’ e-mails	
8	Total successful injections (1 + 3 + 7)	

Table 9: Classification of the e-mails that we received into broad categories of the vulnerability.

We explain the combination of these header injections (4-7) as follows:

- E-Mail Header Injections with “bcc” and “x-check” headers
These represent the perfect attack scenario where we are able to inject multiple headers into the e-mails. We can see that over 75% of the received “bcc” header

injected e-mails are also susceptible to other header injections.

- Both “To” header injections and “x-check” headers
This combination shows us that in addition to being able to inject into the “To” fields, we are able to inject additional headers into the e-mail. It is not clear what causes this behavior; however, these can be exploited to achieve the same result as a regular E-Mail Header Injection.
- “x-check” headers found in ‘nuser’ e-mails
In addition to analyzing the ‘maluser’ account, we also analyze emails received by the ‘nuser’ account. We explain the presence of these headers in the following paragraph.
- Unique “x-check” headers found in ‘nuser’ e-mails
These represent the e-mails with form_ids that were *not* already found in the ‘maluser’ account. We attribute these e-mails to (probably) having a backend that was built with Python or another language having a similar behavior with respect to constructing headers.
- Total successful injections
This represents the total number of successful injections our system made. This includes the E-Mail Header Injections with “bcc” header (1), “To” header injections alone (3), and Unique “x-check” headers found in ‘nuser’ e-mails (7). This is the total number of vulnerabilities that were found by our system.

5.4 Understanding the pipeline

This section serves to represent our pipeline quantitatively and graphically. Table 10 showcases the data gathered by our pipeline, with the differential changes at each stage of the pipeline.

At each stage of the pipeline, the amount of data trickles down, for instance, out of the 21,675,680 URLs we crawled, only 6,794,917 forms (31.35%) were found. Out of these, only 1,132,157 forms (16.66%) contained e-mail fields.

In our fuzzing attempts, the same behavior is repeated. We fuzzed 934,016 forms with the regular payload, which resulted in a total of 52,724 e-mails (5.64%). After analysis of the received e-mails, we further fuzzed 934,016 forms, which resulted in 496 e-mails (1.07%) which contain the vulnerability.

S.No	Pipeline Stage	Quantity	Differential Δ $d2/d1 * 100$
1	Crawled URLs	21,675,680	—
2	Forms found	6,794,917	31.35%
3	E-Mail Forms found	1,132,157	16.66%
4	Fuzzed with regular payload	934,016	82.50%
5	Received e-mails	52,724	5.64%
6	Fuzzed with malicious payload	46,156	87.54%
7	Successful attacks	496	1.07%

Table 10: Data gathered by our pipeline at each stage, with the differential between the stages.

Thus, our pipeline can be visualized as a funnel (shown in Figure 9), where the quantity decreases until it reaches the end product.




Figure 9: Pipeline - shows the quantity of data gathered at each stage of the pipeline.

From our research, it is clear that E-Mail Header Injection is quite widespread as a vulnerability, appearing on 1.07% of forms that we were able to perform automated attacks on. This value acts as a lower bound for E-Mail Header Injection vulnerability, and can quite easily be much more if the attacks were of a more concentrated nature, crafted for the individual websites and less automated. We discuss this possibility, and other concepts such as limitations in the following chapter.

Table 11: Frequency of Special Characters

Non-English or Math	Frequency	Comments
Ø	1 in 1,000	For Swedish names
π	1 in 5	Common in math
\$	4 in 5	Used in business
Ψ ₁ ²	1 in 40,000	Unexplained usage

6. THE BODY OF THE PAPER

Typically, the body of a paper is organized into a hierarchical structure, with numbered or unnumbered headings for sections, subsections, sub-subsections, and even smaller sections. The command `\section` that precedes this paragraph is part of such a hierarchy.⁴ L^AT_EX handles the numbering and placement of these headings for you, when you use the appropriate heading commands around the titles of the headings. If you want a sub-subsection or smaller part to be unnumbered in your output, simply append an asterisk to the command name. Examples of both numbered and unnumbered headings will appear throughout the balance of this sample document.

6.1 Type Changes and Special Characters

We have already seen several typeface changes in this sample. You can indicate italicized words or phrases in your text with the command `\textit`; emboldening with the command `\textbf` and typewriter-style (for instance, for computer code) with `\texttt`. But remember, you do not have to indicate typestyle changes when such changes are part of the *structural* elements of your article; for instance, the heading of this subsection will be in a sans serif⁵ typeface, but that is handled by the document class file. Take care with the use of⁶ the curly braces in typeface changes; they mark the beginning and end of the text that is to be in the different typeface.

You can use whatever symbols, accented characters, or non-English characters you need anywhere in your document; you can find a complete list of what is available in the *L^AT_EX User's Guide*[?].

6.2 Tables

Because tables cannot be split across pages, the best placement for them is typically the top of the page nearest their initial cite. To ensure this proper “floating” placement of tables, use the environment `table` to enclose the table's contents and the table caption. The contents of the table itself must go in the `tabular` environment, to be aligned properly in rows and columns, with the desired horizontal and vertical rules. Again, detailed instructions on `tabular` material is found in the *L^AT_EX User's Guide*.

Immediately following this sentence is the point at which Table 1 is included in the input file; compare the placement of the table here with the table in the printed dvi output of this document.

To set a wider table, which takes up the whole width

⁴This is the second footnote. It starts a series of three footnotes that add nothing informational, but just give an idea of how footnotes work and look. It is a wordy one, just so you see how a longish one plays out.

⁵A third footnote, here. Let's make this a rather short one to see how it looks.

⁶A fourth, and last, footnote.

Figure 10: A sample black and white graphic.

Figure 11: A sample black and white graphic that has been resized with the `includegraphics` command.

of the page's live area, use the environment `table*` to enclose the table's contents and the table caption. As with a single-column table, this wide table will “float” to a location deemed more desirable. Immediately following this sentence is the point at which Table 2 is included in the input file; again, it is instructive to compare the placement of the table here with the table in the printed dvi output of this document.

6.3 Figures

Like tables, figures cannot be split across pages; the best placement for them is typically the top or the bottom of the page nearest their initial cite. To ensure this proper “floating” placement of figures, use the environment `figure` to enclose the figure and its caption.

This sample document contains examples of `.eps` files to be displayable with L^AT_EX. If you work with pdfL^AT_EX, use files in the `.pdf` format. Note that most modern T_EX system will convert `.eps` to `.pdf` for you on the fly. More details on each of these is found in the *Author's Guide*.

As was the case with tables, you may want a figure that spans two columns. To do this, and still to ensure proper “floating” placement of tables, use the environment `figure*` to enclose the figure and its caption. and don't forget to end the environment with `figure*`, not `figure`!

6.4 Theorem-like Constructs

Other common constructs that may occur in your article are the forms for logical constructs like theorems, axioms, corollaries and proofs. There are two forms, one produced by the command `\newtheorem` and the other by the command `\newdef`; perhaps the clearest and easiest way to distinguish them is to compare the two in the output of this sample document:

This uses the `theorem` environment, created by the `\newtheorem` command:

THEOREM 1. *Let f be continuous on $[a,b]$. If G is an antiderivative for f on $[a,b]$, then*

$$\int_a^b f(t)dt = G(b) - G(a).$$

The other uses the `definition` environment, created by the `\newdef` command:

Definition 1. *If z is irrational, then by e^z we mean the unique number which has logarithm z :*

$$\log e^z = z$$

Two lists of constructs that use one of these forms is given in the *Author's Guidelines*.

There is one other similar construct environment, which is already set up for you; i.e. you must *not* use a `\newdef` command to create it: the `proof` environment. Here is a example of its use:

Table 12: Some Typical Commands

Command	A Number	Comments
<code>\alignauthor</code>	100	Author alignment
<code>\numberofauthors</code>	200	Author enumeration
<code>\table</code>	300	For tables
<code>\table*</code>	400	For wider tables

Figure 12: A sample black and white graphic that needs to span two columns of text.

PROOF. Suppose on the contrary there exists a real number L such that

$$\lim_{x \rightarrow \infty} \frac{f(x)}{g(x)} = L.$$

Then

$$l = \lim_{x \rightarrow c} f(x) = \lim_{x \rightarrow c} \left[g(x) \cdot \frac{f(x)}{g(x)} \right] = \lim_{x \rightarrow c} g(x) \cdot \lim_{x \rightarrow c} \frac{f(x)}{g(x)} = 0 \cdot L = 0,$$

which contradicts our assumption that $l \neq 0$. \square

Complete rules about using these environments and using the two different creation commands are in the *Author's Guide*; please consult it for more detailed instructions. If you need to use another construct, not listed therein, which you want to have the same formatting as the Theorem or the Definition[?] shown above, use the `\newtheorem` or the `\newdef` command, respectively, to create it.

A Caveat for the T_EX Expert

Because you have just been given permission to use the `\newdef` command to create a new form, you might think you can use T_EX's `\def` to create a new command: *Please refrain from doing this!* Remember that your L^AT_EX source code is primarily intended to create camera-ready copy, but may be converted to other forms – e.g. HTML. If you inadvertently omit some or all of the `\defs` recompilation will be, to say the least, problematic.

7. CONCLUSIONS

This paragraph will end the body of this sample document. Remember that you might still have Acknowledgments or Appendices; brief samples of these follow. There is still the Bibliography to deal with; and we will make a disclaimer about that here: with the exception of the reference to the L^AT_EX book, the citations in this paper are to articles which have nothing to do with the present subject and are used as examples only.

8. ACKNOWLEDGMENTS

This section is optional; it is a location for you to acknowledge grants, funding, editing assistance and what have you. In the present case, for example, the authors would like to thank Gerald Murray of ACM for his help in codifying this *Author's Guide* and the `.cls` and `.tex` files that it describes.

References

- [1] T. Berners-Lee, L. Masinter, and M. McCahill. Internet Message Format - RFC 1738. 2008.

- [2] S. W. Boyd and A. D. Keromytis. Sqlrand: Preventing sql injection attacks. In *Applied Cryptography and Network Security*, pages 292–302. Springer, 2004.
- [3] Celery. Celery Homepage, 2016.
- [4] 2010.
- [5] W. G. Halfond, J. Viegas, and A. Orso. A classification of sql-injection attacks and countermeasures. In *Proceedings of the IEEE International Symposium on Secure Software Engineering*, volume 1, pages 13–15. IEEE, 2006.
- [6] A. Herzog. Full Disclosure: JavaMail SMTP Header Injection via method setSubject [CSNC-2014-001], 2014.
- [7] H. H. Injection. HTTP header injection — Wikipedia, the free encyclopedia. <http://en.wikipedia.org/w/index.php?title=HTTP%20header%20injection&oldid=713295668>, 2016. [Online; accessed 18-April-2016].
- [8] T. Jim, N. Swamy, and M. Hicks. Defeating script injection attacks with browser-enforced embedded policies. In *Proceedings of the 16th International Conference on World Wide Web, WWW '07*, pages 601–610, New York, NY, USA, 2007. ACM.
- [9] M. Johns and J. Winter. Requestrodeo: Client side protection against session riding. In *Proceedings of the OWASP Europe 2006 Conference*, 2006.
- [10] A. Klein. [DOM Based Cross Site Scripting or XSS of the Third Kind] Web Security Articles - Web Application Security Consortium, 2005.
- [11] D. Kohler. damonkohler.com: Email Injection, 2008.
- [12] D. McGrath. HTTP requests optimization, 2009.
- [13] OWASP. OWASP Top Ten Project, 2013.
- [14] Phishing. Phishing — Wikipedia, the free encyclopedia. <http://en.wikipedia.org/w/index.php?title=Phishing&oldid=706223617>, 2016. [Online; accessed 27-February-2016].
- [15] PHP-Manual. PHP mail - Send mail, 2016.
- [16] C. Pope and K. Kaur. Is it human or computer? defending e-commerce with captchas. *IT Professional*, 7(2):43–49, Mar 2005.

- [17] A. Sadeghian, M. Zamani, and A. A. Manaf. A taxonomy of sql injection detection and prevention techniques. In *Informatics and Creative Multimedia (ICICM), 2013 International Conference on*, pages 53–56. IEEE, 2013.
- [18] V. Shkapenyuk Torsten Suel. Design and Implementation of a High-Performance Distributed Web Crawler. 2001.
- [19] D. Stuttard and M. Pinto. *The Web Application Hacker’s Handbook: Finding and Exploiting Security Flaws*. John Wiley & Sons, 2011.
- [20] Tobozo. Mail headers injections with PHP, 2004.
- [21] L. von Ahn, M. Blum, and J. Langford. Telling humans and computers apart automatically. *Commun. ACM*, 47(2):56–60, Feb. 2004.
- [22] W3techs. Usage Statistics and Market Share of PHP for Websites, February 2016, 2016.
- [23] Wikipedia. Black-box testing — Wikipedia, the free encyclopedia. <http://en.wikipedia.org/w/index.php?title=Black-box%20testing&oldid=702083755>, 2016. [Online; accessed 02-March-2016].
- [24] Wikipedia. Law of averages — Wikipedia, the free encyclopedia. <http://en.wikipedia.org/w/index.php?title=Law%20of%20averages&oldid=706716293>, 2016. [Online; accessed 08-March-2016].
- [25] Wikipedia. Law of large numbers — Wikipedia, the free encyclopedia. <http://en.wikipedia.org/w/index.php?title=Law%20of%20large%20numbers&oldid=706596753>, 2016. [Online; accessed 08-March-2016].
- [26] Wikipedia. RabbitMQ — Wikipedia, the free encyclopedia. <http://en.wikipedia.org/w/index.php?title=RabbitMQ&oldid=708777848>, 2016. [Online; accessed 09-March-2016].
- [27] Wikipedia. Regression toward the mean — Wikipedia, the free encyclopedia. <http://en.wikipedia.org/w/index.php?title=Regression%20toward%20the%20mean&oldid=703369877>, 2016. [Online; accessed 08-March-2016].
- [28] Wikipedia. Spoofing attack — Wikipedia, the free encyclopedia. <http://en.wikipedia.org/w/index.php?title=Spoofing%20attack&oldid=711407565>, 2016. [Online; accessed 23-April-2016].

APPENDIX

A. HEADINGS IN APPENDICES

The rules about hierarchical headings discussed above for the body of the article are different in the appendices. In the **appendix** environment, the command **section** is used to indicate the start of each Appendix, with alphabetic order designation (i.e. the first is A, the second B, etc.) and a title (if you include one). So, if you need hierarchical structure *within* an Appendix, start with **subsection** as the highest level. Here is an outline of the body of this document in Appendix-appropriate form:

A.1 Introduction

A.2 The Body of the Paper

A.2.1 Type Changes and Special Characters

A.2.2 Math Equations

Inline (In-text) Equations.

Display Equations.

A.2.3 Citations

A.2.4 Tables

A.2.5 Figures

A.2.6 Theorem-like Constructs

A Caveat for the T_EX Expert

A.3 Conclusions

A.4 Acknowledgments

A.5 Additional Authors

This section is inserted by L^AT_EX; you do not insert it. You just add the names and information in the `\additionalauthors` command at the start of the document.

A.6 References

Generated by bibtex from your .bib file. Run latex, then bibtex, then latex twice (to resolve references) to create the .bbl file. Insert that .bbl file into the .tex source file and comment out the command `\thebibliography`.

B. MORE HELP FOR THE HARDY

The sig-alternate.cls file itself is chock-full of succinct and helpful comments. If you consider yourself a moderately experienced to expert user of L^AT_EX, you may find reading it useful but please remember not to change it.